

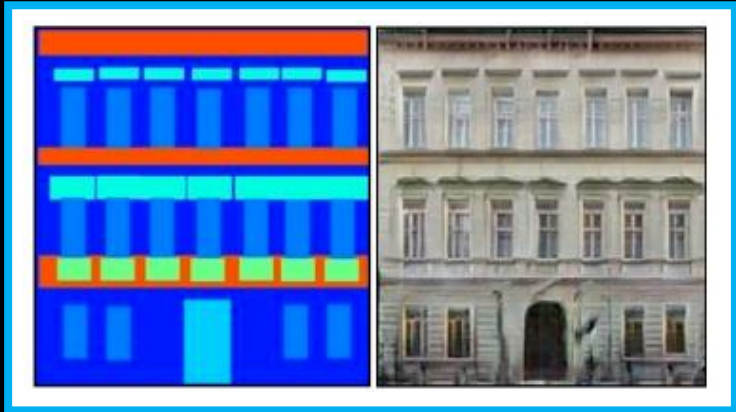
AI for the Media

Week 9, Domain 2 Domain



Motivation for today

Overview of additional machine learning techniques:

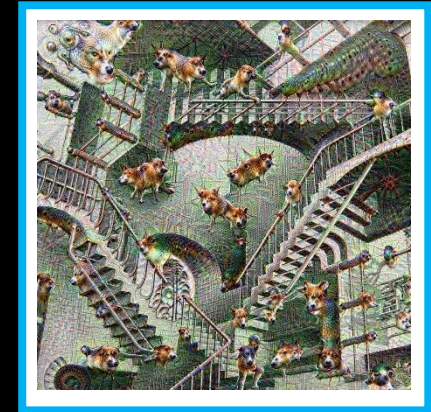


pix2pix

& co



style transfer



deep dream

*& adversarial
attacks*

Overview


Domain 2 Domain Models (*pre-recorded lecture*):

- **Pix2pix** and domain to domain transfer
- **Style transfer** technique
- **Deep Dream** technique


Practical session (*during the live session*):

- **Code**: Training a pix2pix model

Domain to Domain


- We have seen few models working with images on outputs and inputs (for example [AutoEncoders](#)). Similar architectures can be also used to model **relation between two domains** of data. 

Domain to Domain


- We have seen few models working with images on outputs and inputs (for example [AutoEncoders](#)). Similar architectures can be also used to model **relation between two domains** of data. 

- Let's look at a predecessor of this idea, the **U-Net model**:
 - Paper with U-Net applied on the task of lunar crater identification


(PS: U-Net is from 2015, this [paper](#) from 2018)





U-Net



U-Net




Images from one
“domain”





Paired images from
another “domain”

- Connects one image to another while going through a bottleneck

U-Net



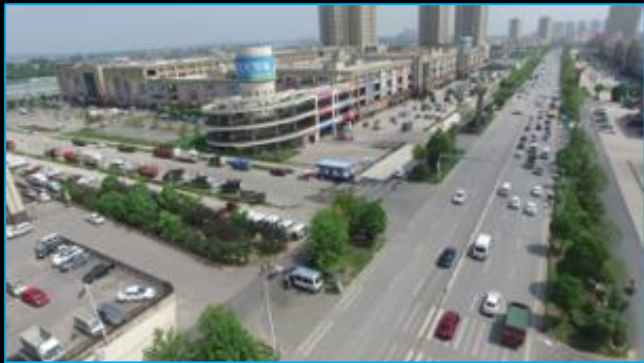
Images from one
“domain”





Paired images from
another “domain”

- Additionally: “*skip connections*” that allow the information to flow

U-Net




Images from one
“domain”




(if this was an
autoencoder)

- PS: As an autoencoder, this model could very easily cheat ...

U-Net in a real-world scenario




< Slightly more experimental setup
– check Siamese U-Nets*




- Can be adapted to many tasks – for example this “change detection”, where it needed to learn *which kind of change* we care about

Pix2Pix


- General purpose **image-to-image translation**:
 - From their paper: “Many problems in image processing, computer graphics, and computer vision can be posed as “*translating*” an input image into a corresponding output image.”




Pix2Pix



- General purpose **image-to-image translation**:
 - From their paper: “Many problems in image processing, computer graphics, and computer vision can be posed as “*translating*” an input image into a *corresponding output image*.”
 - Translating without specifically defining the rules – data-driven translating between two domains by showing **paired** examples:
 - [Image from A, Corresponding Image from B] * N samples



Pix2Pix



- General purpose **image-to-image translation**:

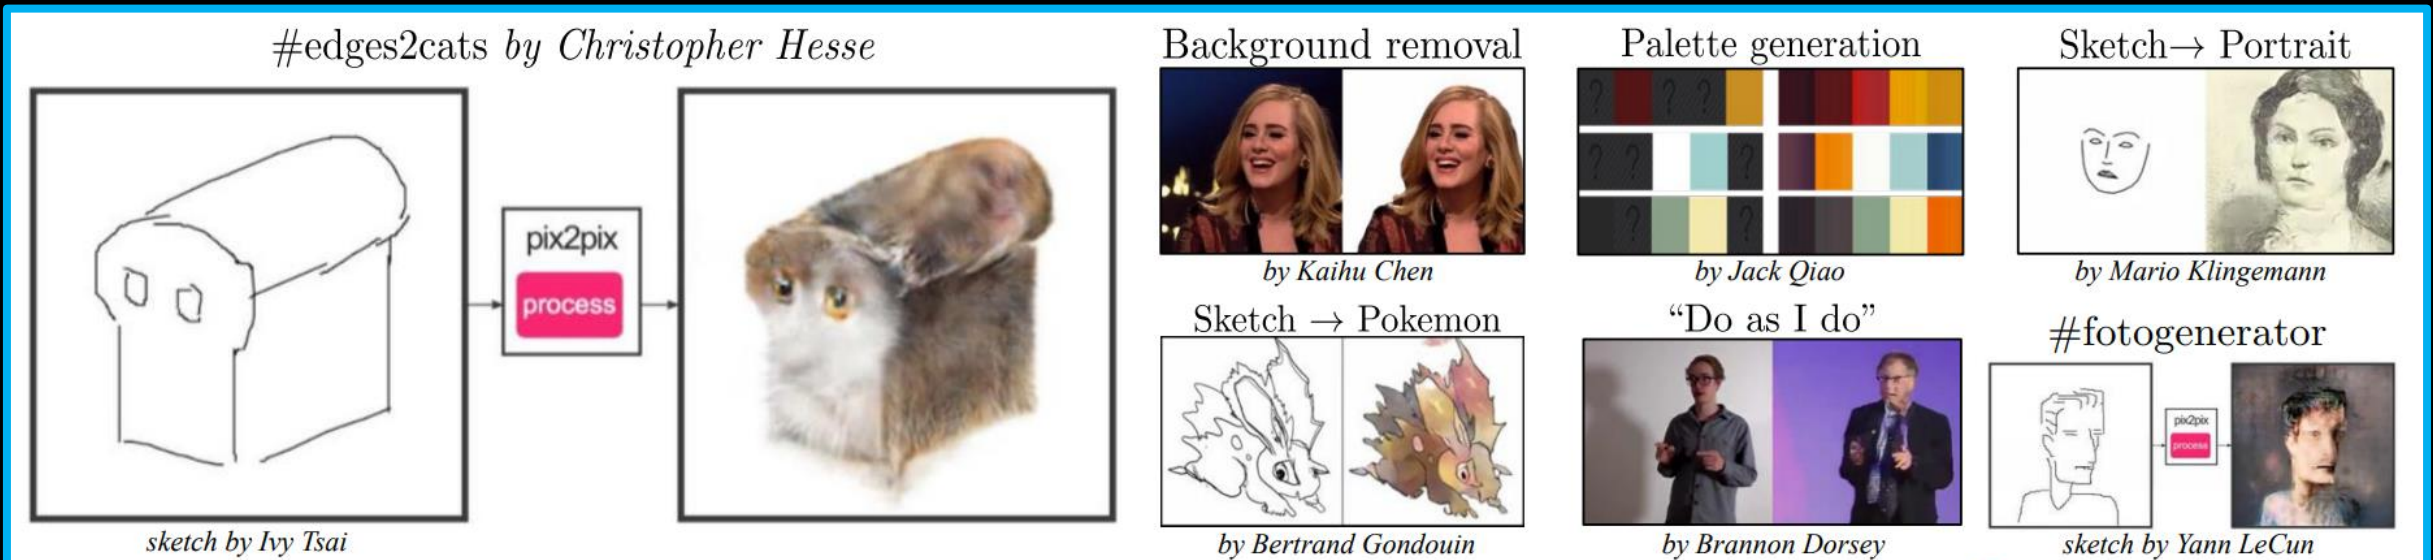
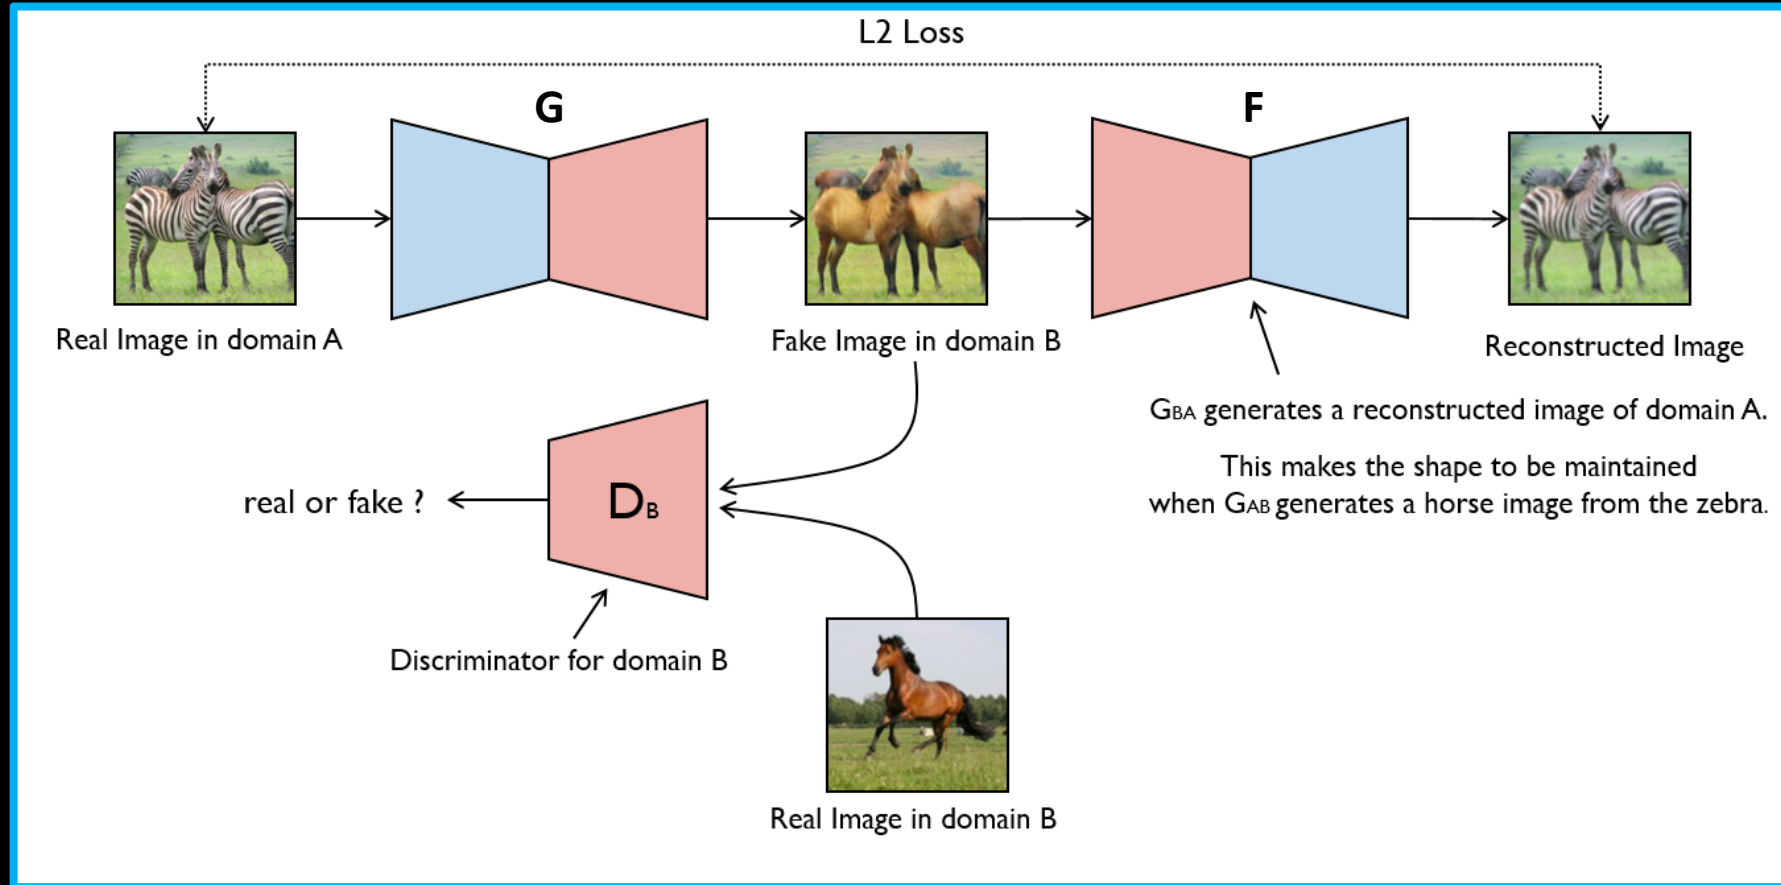


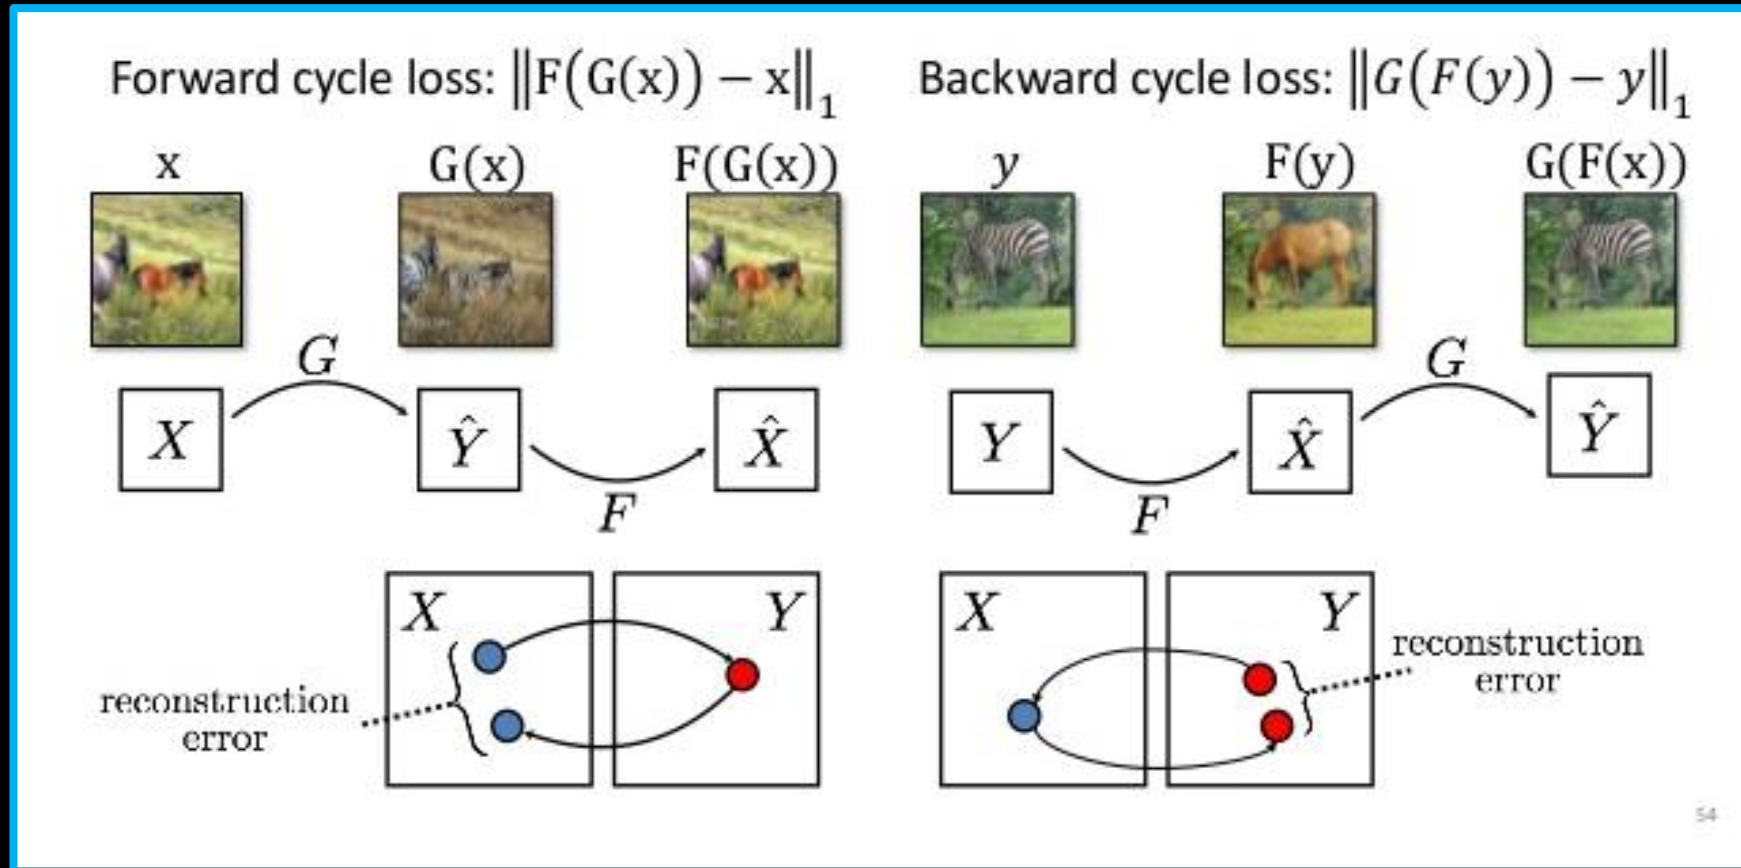
Figure 11: Example applications developed by online community based on our `pix2pix` codebase: `#edges2cats` [3] by Christopher Hesse, *Background removal* [6] by Kaihu Chen, *Palette generation* [5] by Jack Qiao, *Sketch → Portrait* [7] by Mario Klingemann, *Sketch → Pokemon* [1] by Bertrand Gondouin, “Do As I Do” pose transfer [2] by Brannon Dorsey, and `#fotogenerator` by Bosman et al. [4].

CycleGAN, connectivity



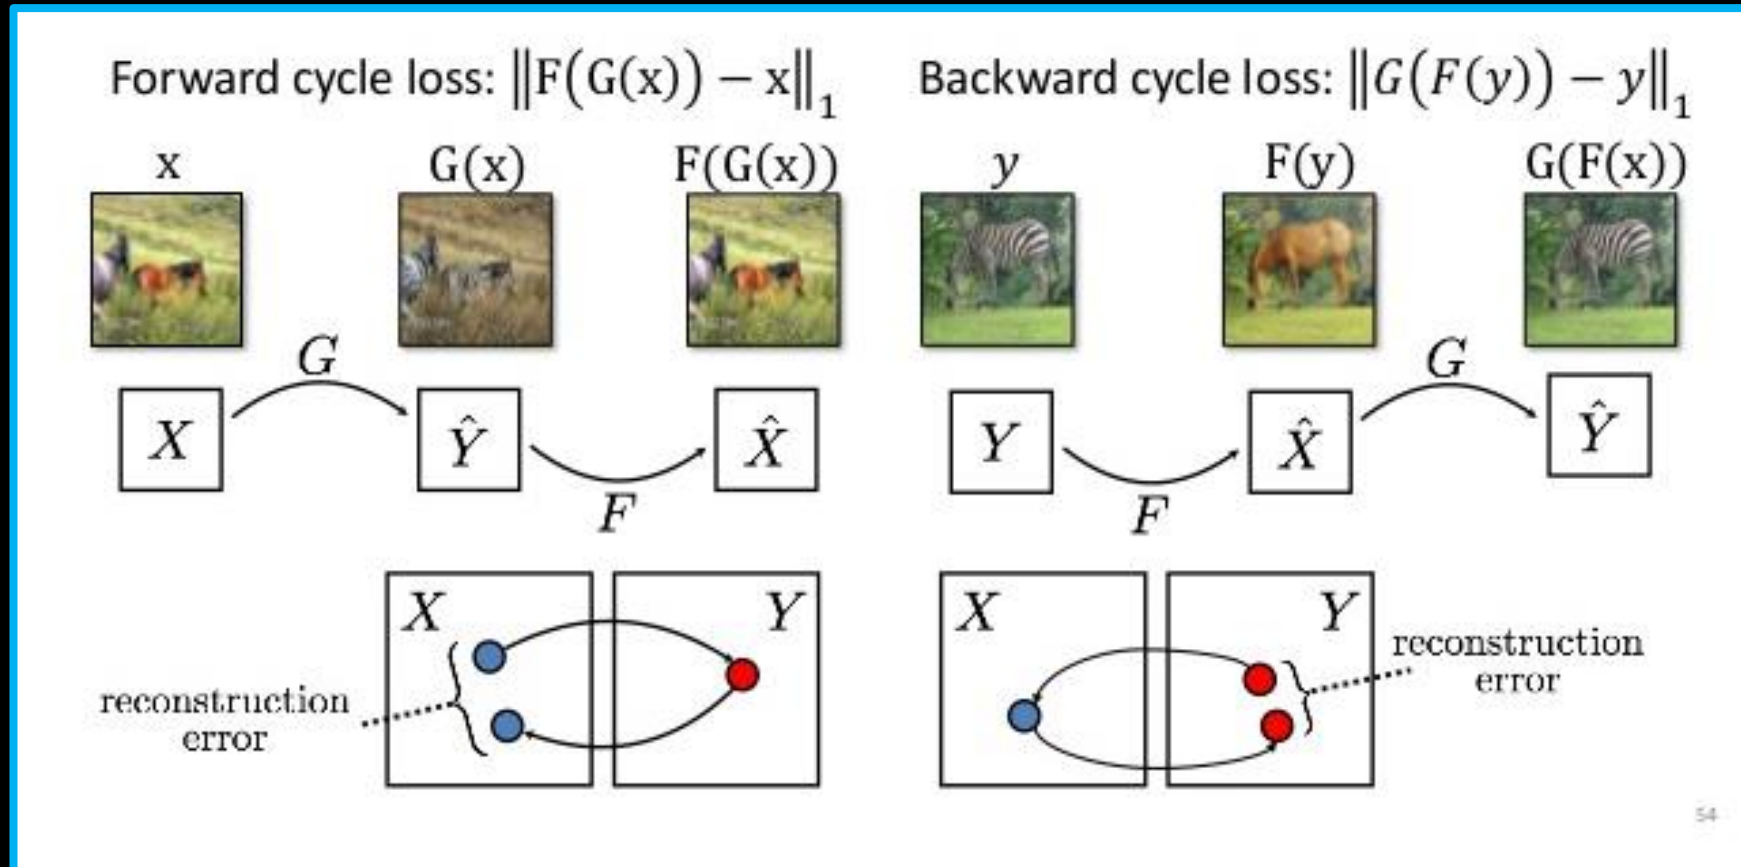
- Setup with two generators going between two domains (G and F) and a special cyclic loss function. This model **doesn't require paired data** in the training datasets.

CycleGAN, loss functions



- And a **cyclic loss function** which looks at two types of errors ^

CycleGAN, loss functions




Check
this
[video](#)


- And a **cyclic loss function** which looks at two types of errors ^

CycleGAN examples



Input



Output



real nighttime



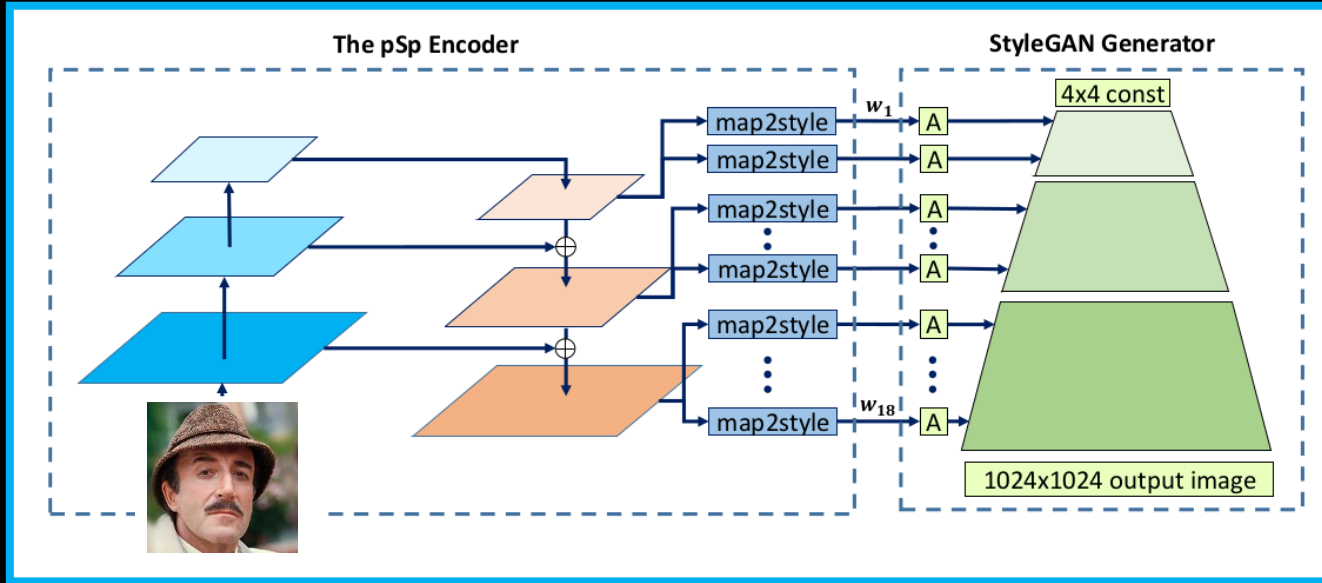
ToDayGAN
(ours)



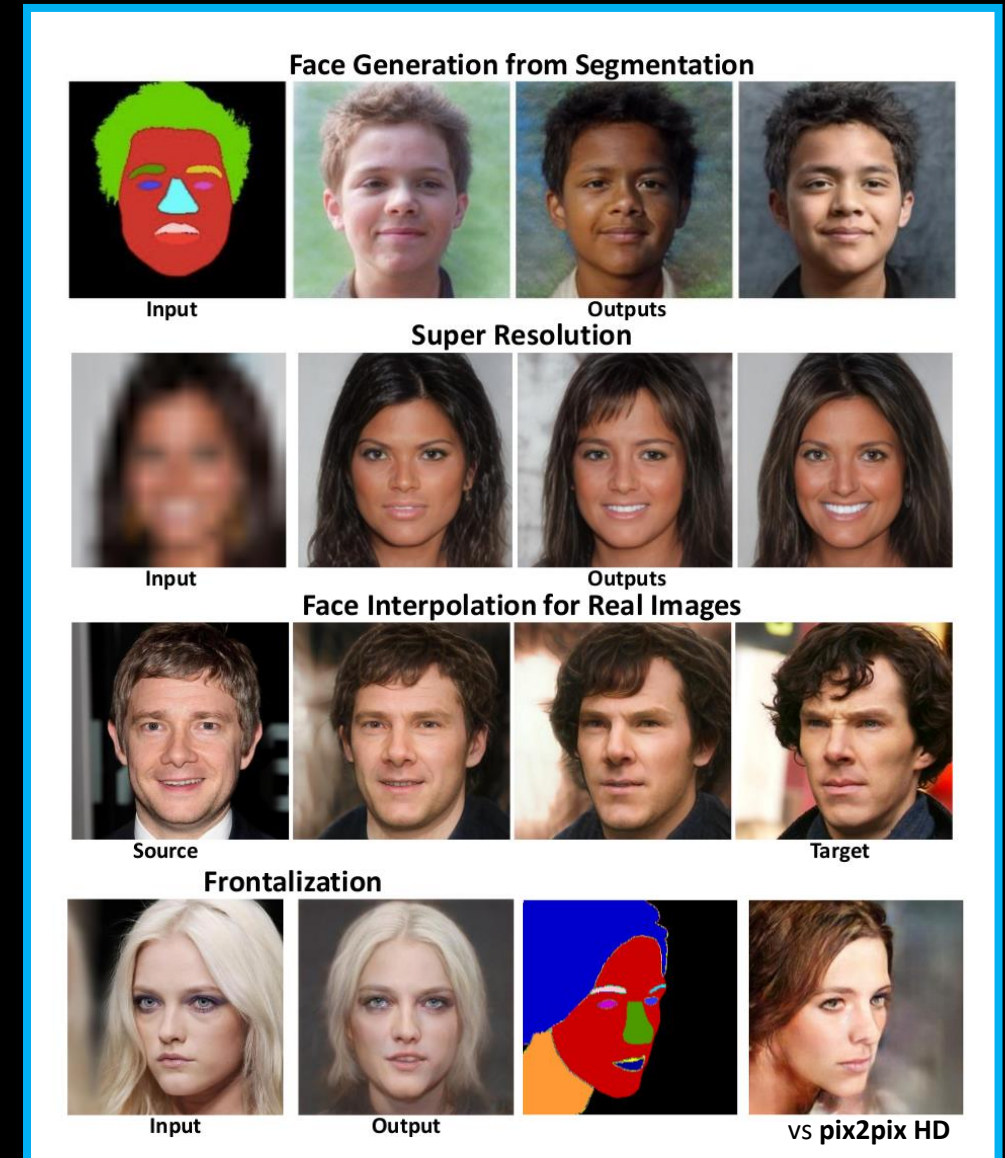
Monet 2 Real, and more examples
at junyanz.github.io/CycleGAN/

Night 2 Day, followed by more tasks

Follow-up works



- **Pixel2Style2Pixel**, new research (2020) which uses a frozen pre-trained StyleGAN2 generator and **trains an encoder** for many specialized tasks while getting much higher quality than a pix2pix would →



Style Transfer

- Previously we saw, how a Deep Convolutional network separates where it saves *high-level* and *low-level* representations (this is due to the *Conv->Pool* combo – each convolutional layer serves as an image-filter).

Style Transfer

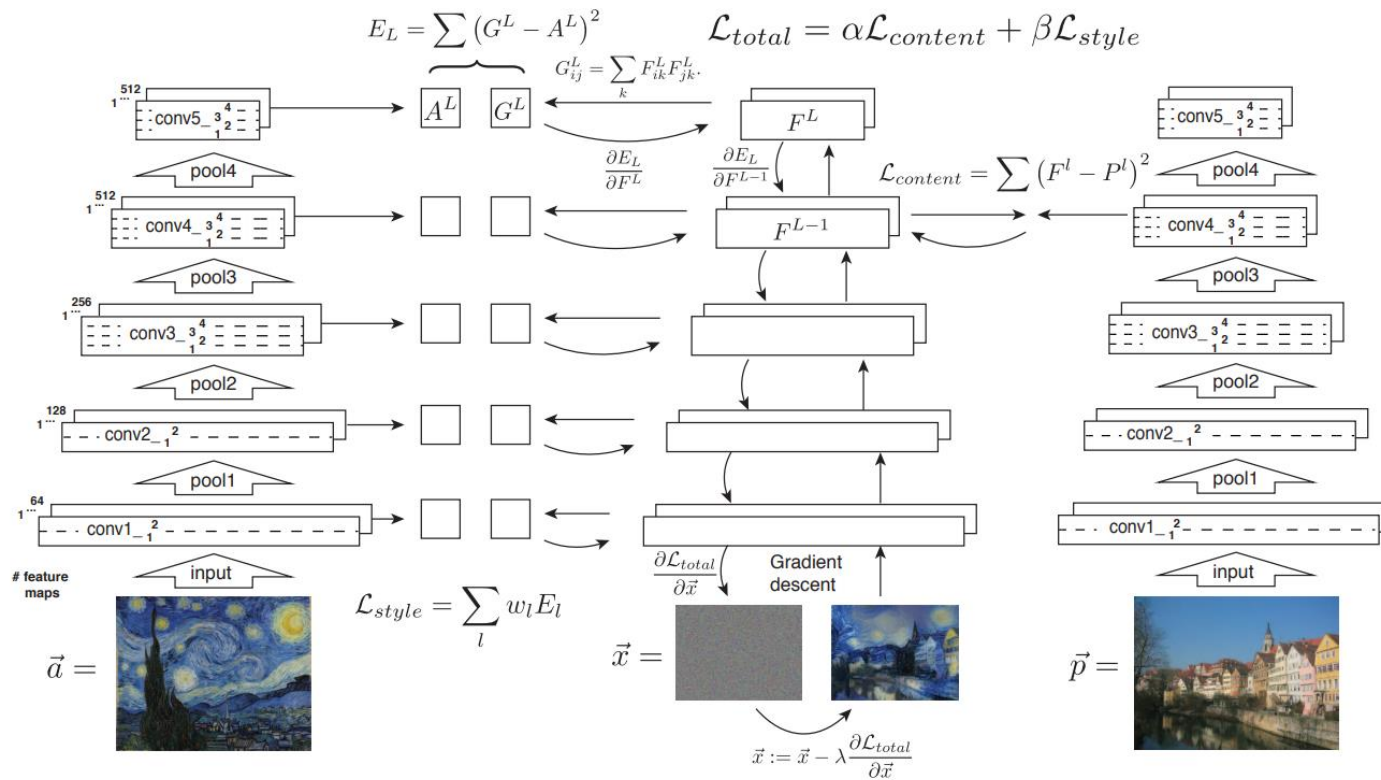
- Previously we saw, how a Deep Convolutional network separates where it saves *high-level* and *low-level* representations (this is due to the *Conv->Pool* combo – each convolutional layer serves as an image-filter).
 - **High-level information** – contents of the whole image – **image content**
 - Encode one image extracting its *content information*, the feature responses in deeper layers of the network

Style Transfer

- Previously we saw, how a Deep Convolutional network separates where it saves *high-level* and *low-level* representations (this is due to the *Conv->Pool* combo – each convolutional layer serves as an image-filter).
 - **High-level information** – contents of the whole image – **image content**
 - Encode one image extracting its *content information*, the feature responses in deeper layers of the network
 - **Low-level information** – details used inside the image, texture – **image style**
 - Encode another image extracting the *style information*, feature response alongside a selection of layers

(PS: One possible similarity if the content of low and high frequencies in music converted to spectrograms)

Style Transfer



Style Transfer

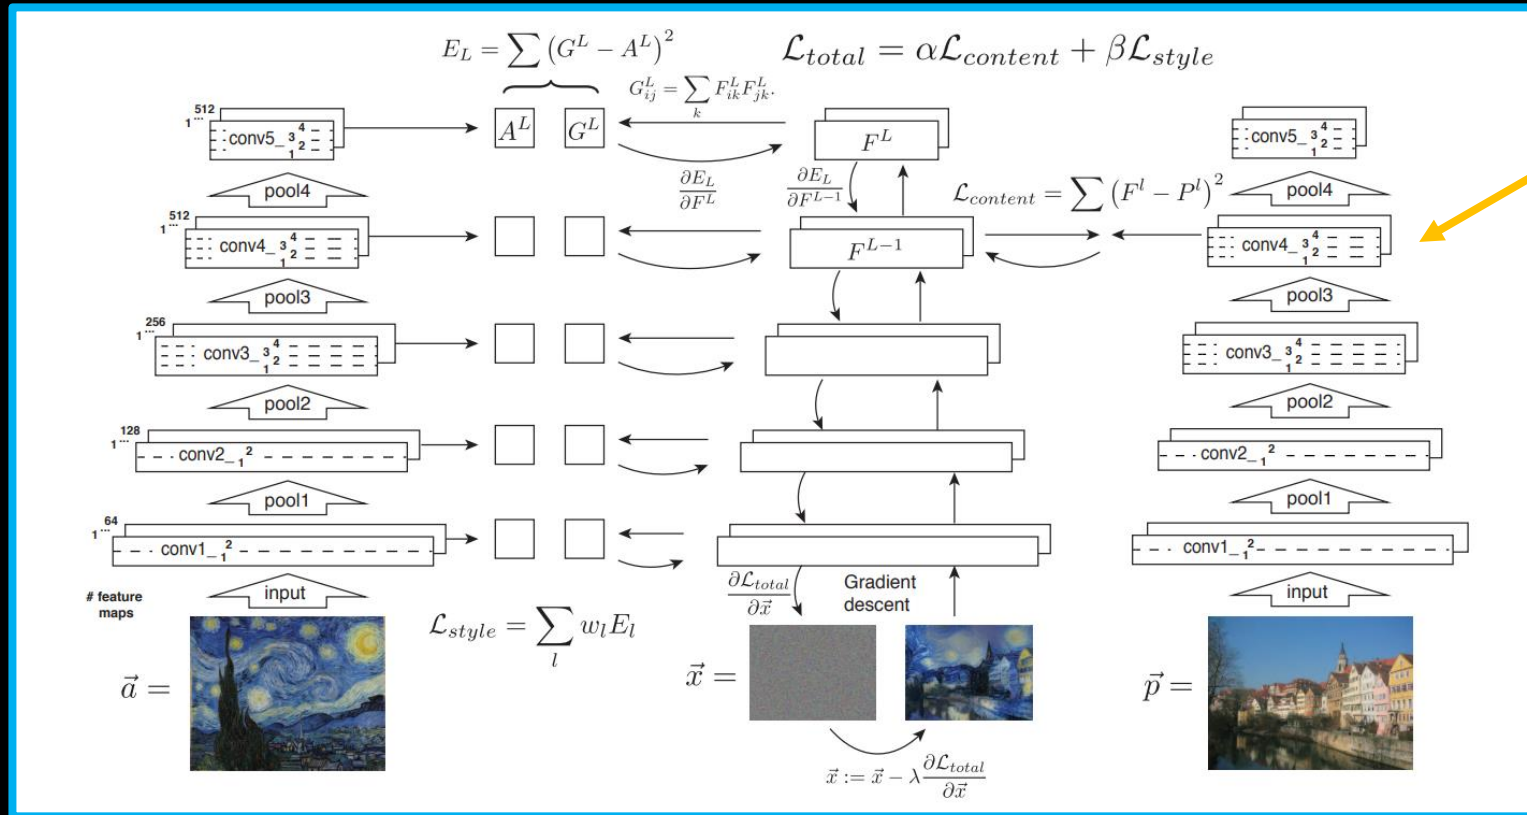


Image content =
the feature
responses in deeper
layers of the VGG
network

Style Transfer

Image style =
feature response
alongside a
selection of layers

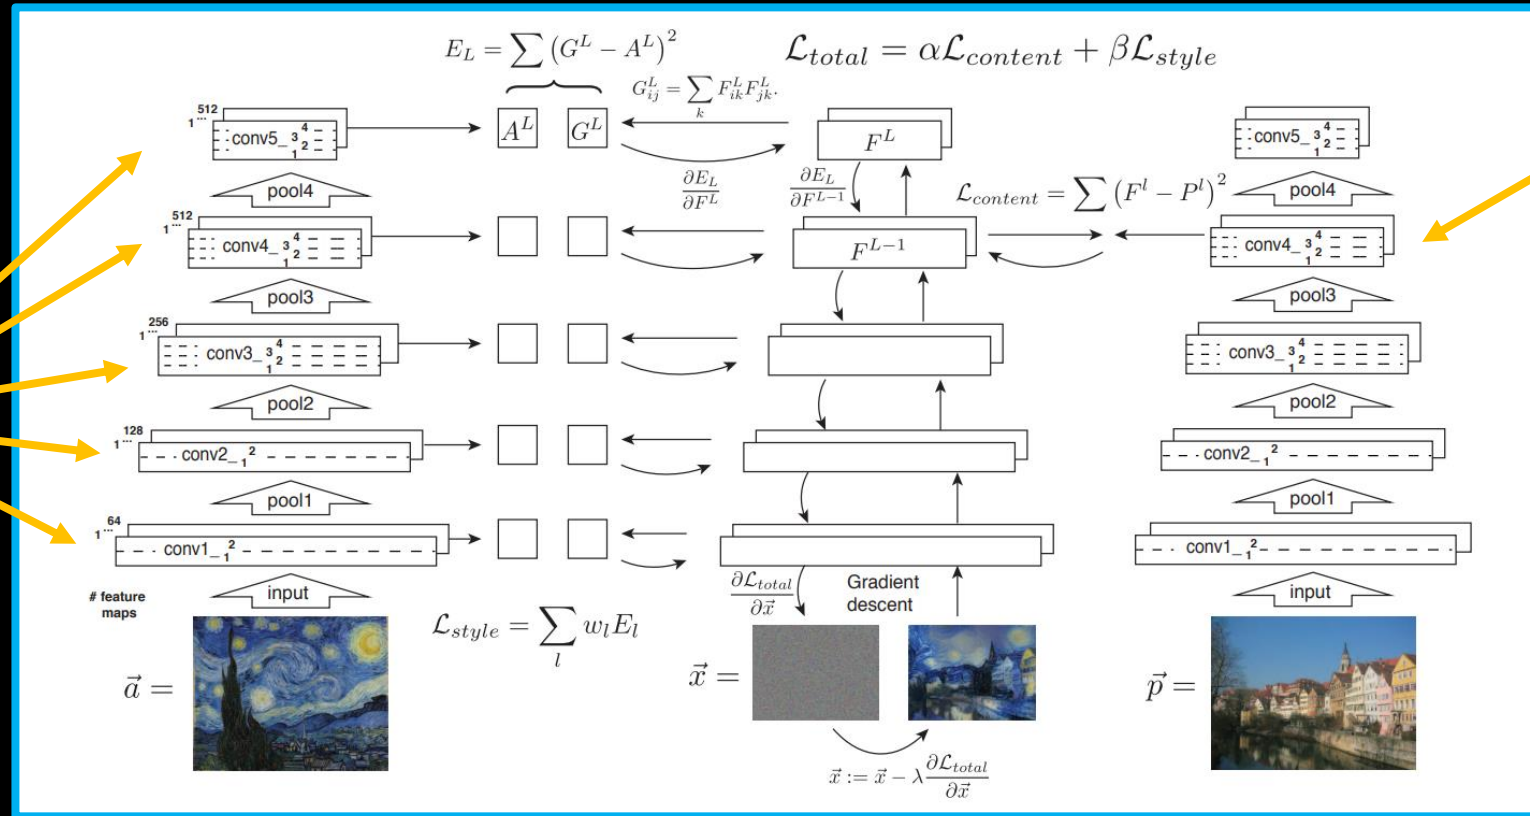


Image content =
the feature
responses in deeper
layers of the VGG
network

Style Transfer

Image style =
feature response
alongside a
selection of layers

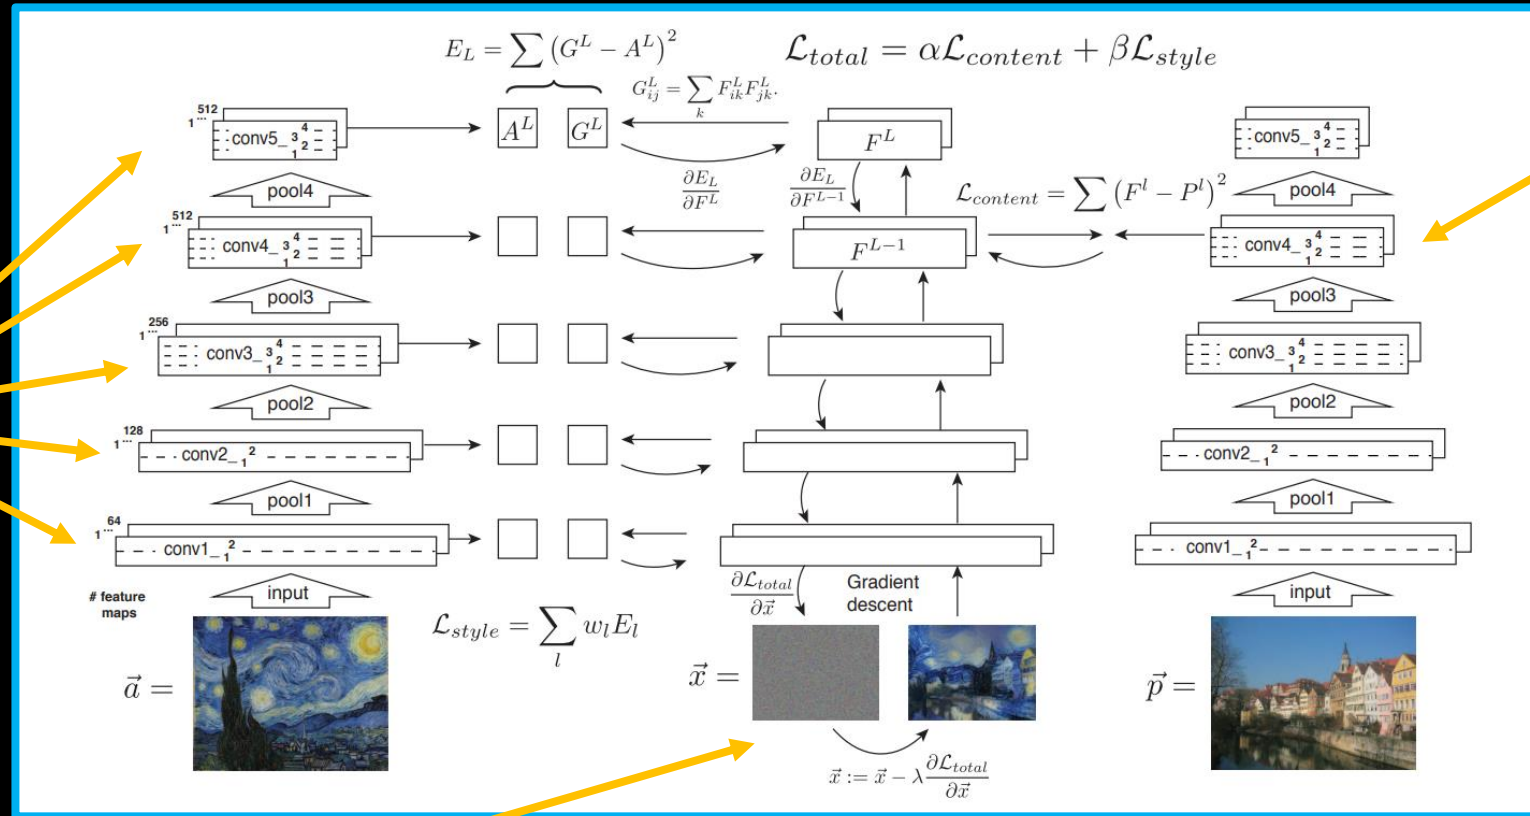


Image content =
the feature
responses in deeper
layers of the VGG
network

Optimizing random noise image to have the same responses as those features we saved.

- Iteratively we will create a new image which has the style responses similar to our encoded style features + content responses similar to our encoded content features

Style Transfer

Image style =
feature response
alongside a
selection of layers

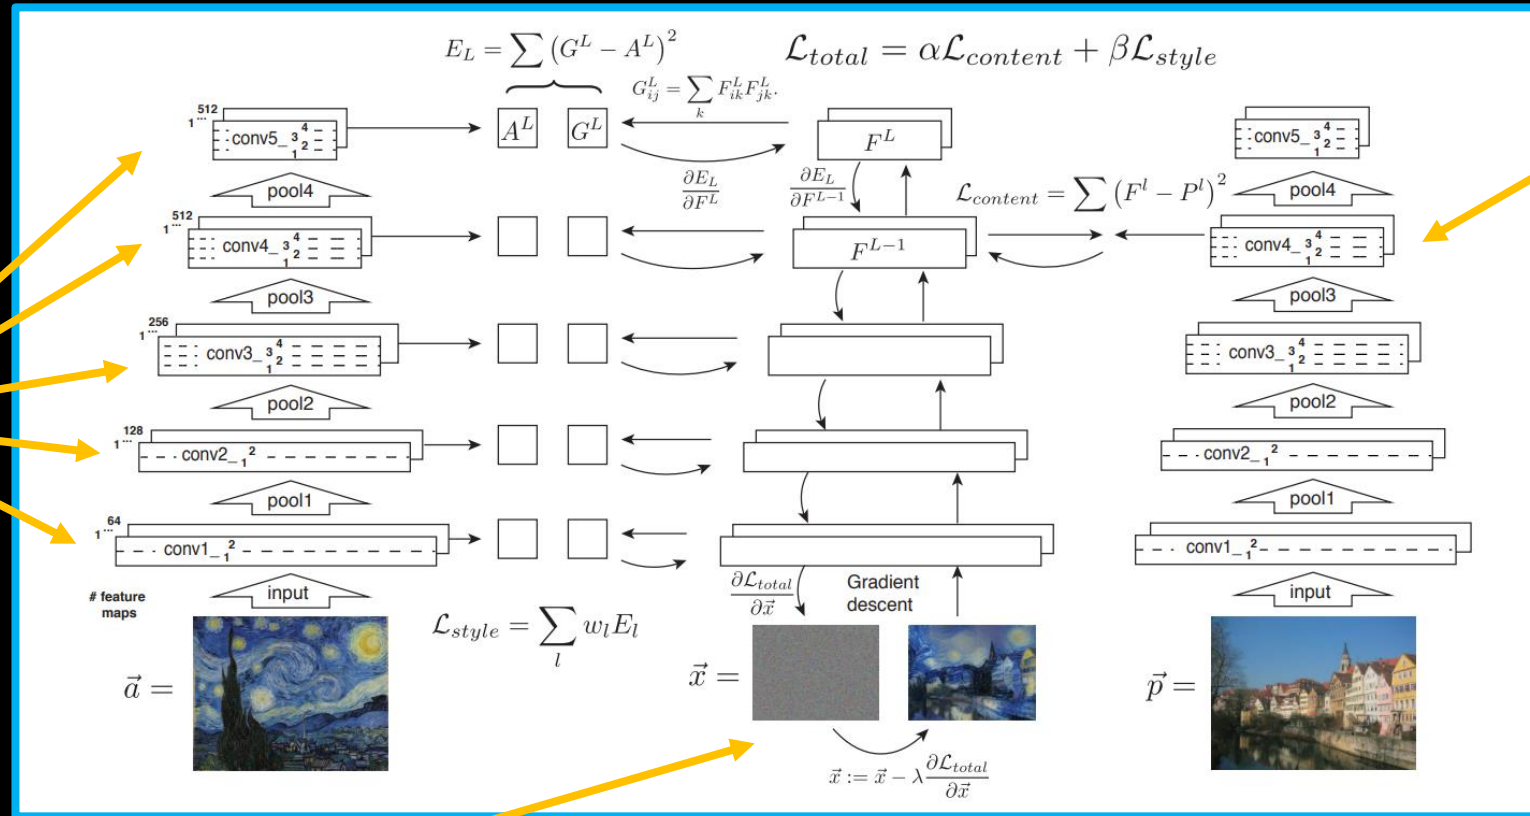


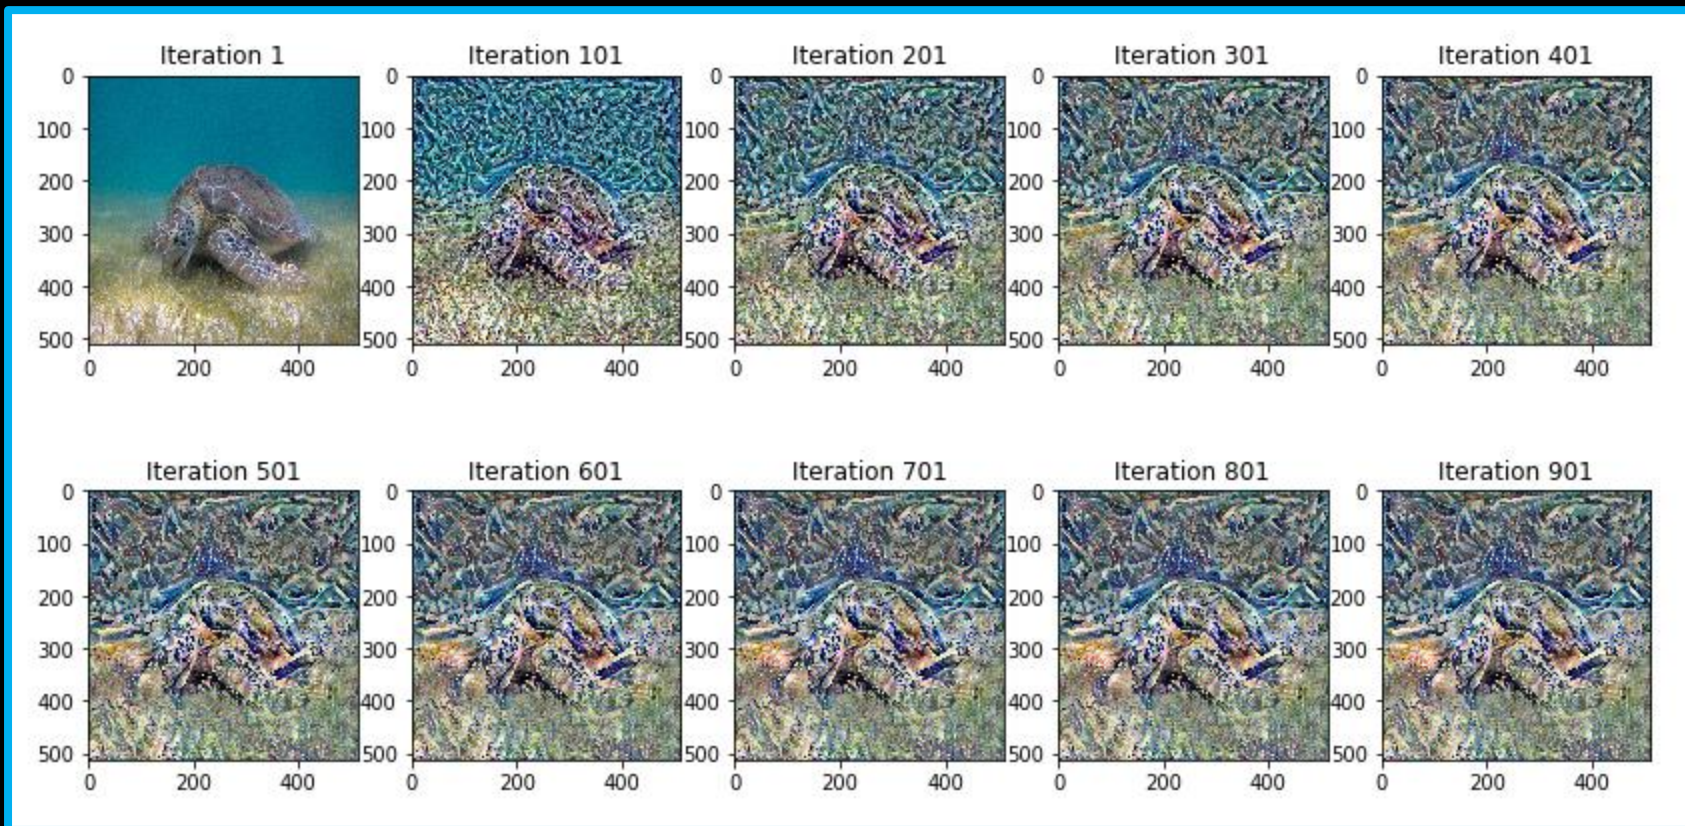
Image content =
the feature
responses in deeper
layers of the VGG
network

Optimizing random noise image to have the same responses as those features we saved.

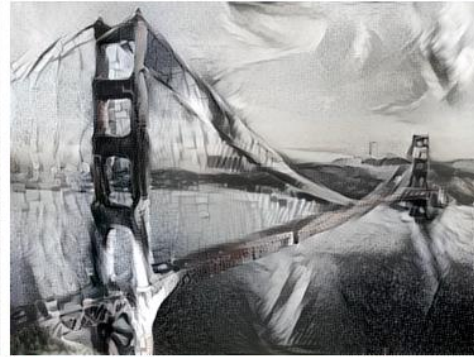
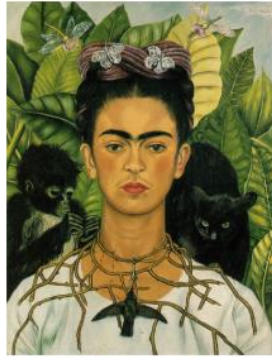
- Iteratively we will create a new image which has the style responses similar to our encoded style features + content responses similar to our encoded content features

This is **relatively slow** (iterative optimization of the input image), later papers sped it up by using an **image2image translation method** with feed-forward networks.

Style Transfer



Style Transfer




Style Transfer

- **Online demos:** deepart.io/latest/
- **Colab notebooks:**
 - Basic style transfer (with arbitrary images): ArtML / [style transfer keras.ipynb](#)
 - Fast style transfer (with pretrained styles): ArtML / [fast-style-transfer](#)
- **Papers:** [style transfer \(2015\)](#), [fast style transfer \(2016\)](#)

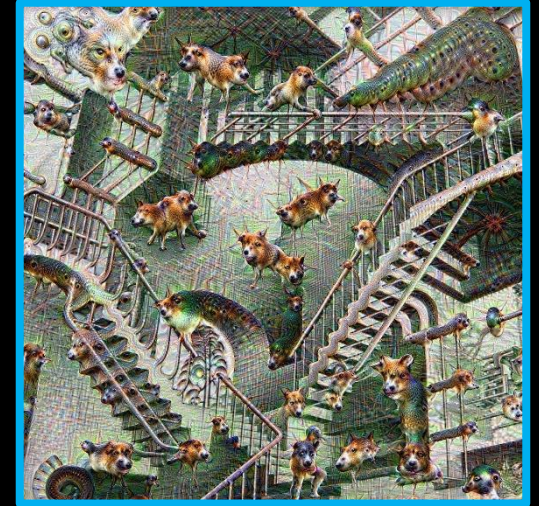
Deep Dream

- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).

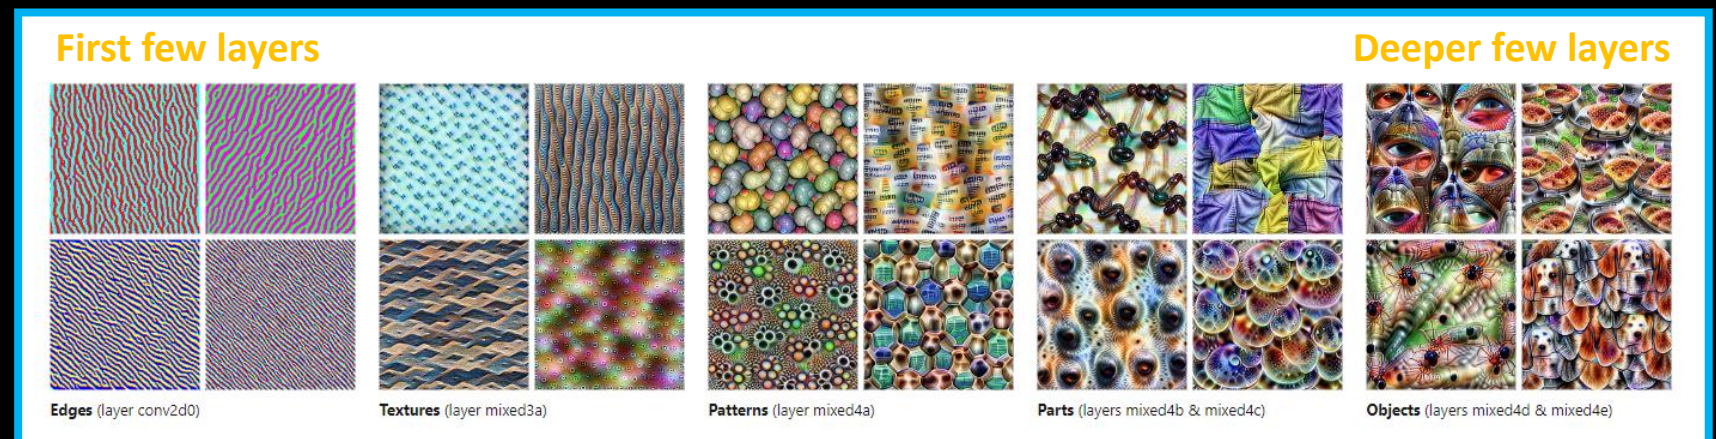


Deep Dream

- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).




Using **Convolutional network** (GoogLeNet) trained for classification on ImageNet:



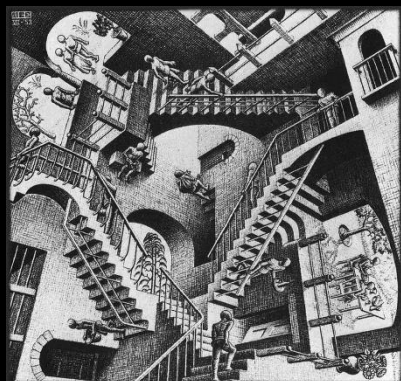
(PS: image [source](#), features [viz.](#), illustrational [video](#))

Deep Dream

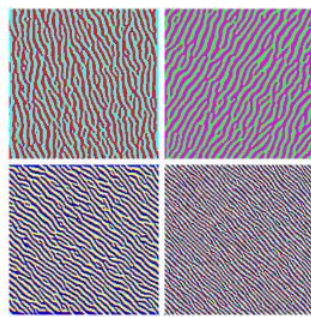
- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).



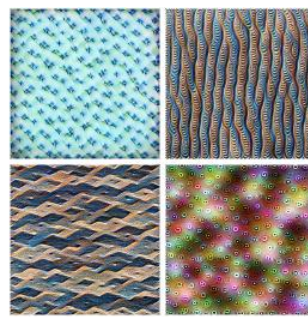
Change this **original image** so that it activates a **selected feature** with the highest possible force!



First few layers



Edges (layer conv2d0)



Textures (layer mixed3a)



Patterns (layer mixed4a)



Parts (layers mixed4b & mixed4c)

Deeper few layers

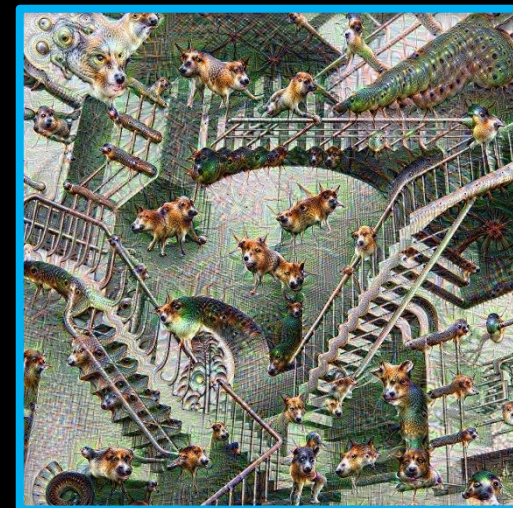


Objects (layers mixed4d & mixed4e)

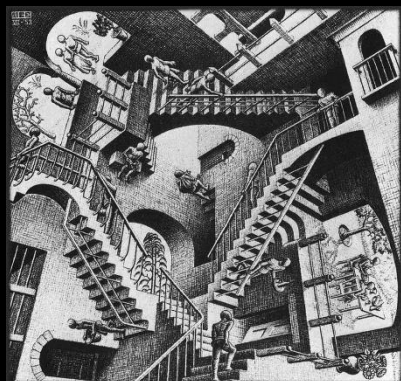
(PS: image [source](#), features [viz.](#), illustrational [video](#))

Deep Dream

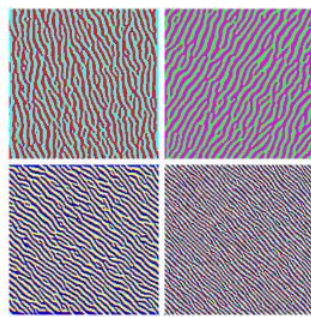
- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).




Change this **original image** so that it activates a **selected feature** with the highest possible force!



First few layers



Edges (layer conv2d0)



Textures (layer mixed3a)



Patterns (layer mixed4a)



Parts (layers mixed4b & mixed4c)

Deeper few layers



Objects (layers mixed4d & mixed4e)



selected feature

(PS: image [source](#), features [viz.](#), illustrational [video](#))

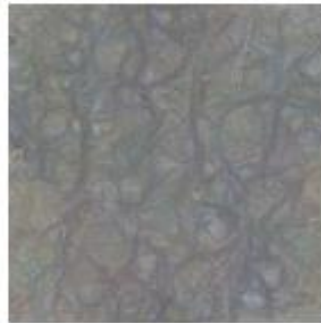
Deep Dream

Iterations:

Starting from random noise, we optimize an image to activate a particular neuron (layer mixed4a, unit 11).



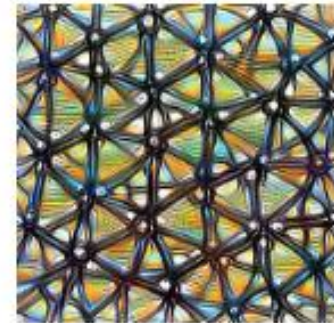
Step 0



Step 4



Step 48



Step 2048

Deep Dream

As a network visualization technique ...

Deep dreaming can reveal information stored inside the network ...
Recall shapes and content information from the originally used dataset.

- Model released by Yahoo to identify NSFW content (without releasing the sensitive information about the dataset): github.com/yahoo/open_nsfw
- Which was (*of course* ...) soon followed by using a deep-dream–like technique: open_nsfw.gitlab.io/
 - Interesting usage of the network to generate *suggestive imagery* (changing shapes of landscapes imagery, etc.)



Deep Dream

- **Online demos:** dreamscopeapp.com *(not 100% sure if it's not a style transfer of deep dream like effect)*
- **Colab notebooks:**
 - Deep dream a photo: [as a ML4A guide](#)
 - Alternative code: ArtML / [neural-synth-clustering-v2.ipynb](#)
- **Reading:** distill.pub/2017/feature-visualization/

Links and additional readings:

- **Online pix2pix demos:** affinelayer.com/pixsrv/
 - **Follow-up papers:** [pix2pixHD](#) (with high. res.), [vid2vid](#) (with frame to frame consistency)
- **Bonus readings:**
 - **Feature Visualization**, Distill – [blog](#)
 - **Sensory Optimization:** Neural Networks as a Model for Understanding and Creating Art – [paper](#)
 - About **Pix2Pix** on ML4A – [blog with code](#)

End of the lecture

*) PS: follows material for the practical session ...

AI for the Media

Week 9, Domain 2 Domain



Practical: pix2pix training

The end